

Quantitative Research Methods

Chapter 7: Moving Beyond Linearity

Prof. Dr. Christoph Weisser

The course at a glance

Chapter	Topic
Ch. 1	Introduction
Ch. 2	Statistical learning — accuracy, bias–variance
Ch. 3	Linear regression — estimation, inference, diagnostics
Ch. 4	Classification — logistic regression, confusion matrix, ROC
Ch. 5	Resampling — validation, CV, bootstrap
Ch. 6	Model selection and regularisation — ridge, lasso
Ch. 7	Moving beyond linearity — splines, GAMs
Ch. 8	Tree-based methods — trees, forests, boosting

Contents

- 1 Why this chapter matters
- 2 Polynomials
- 3 Splines
- 4 Smoothing
- 5 GAMs
- 6 Python Lab
- 7 Summary

Optional and advanced material is collected in the **appendix** at the end of the deck.

Notation in this chapter

Symbol	Meaning
$b_j(x)$	Basis function: fixed transformation of x entering linearly
d	Degree of a polynomial (or of each spline piece)
$C_k(x)$	Indicator (dummy) for the k -th bin of a step function
ξ_k	Knot: breakpoint where adjacent spline pieces join
$(x - \xi)_+^3$	Truncated power term: 0 for $x \leq \xi$, else $(x - \xi)^3$
K	Number of knots; cubic-spline $\text{df} = K + 4$, natural-spline $\text{df} = K$
λ	Smoothing parameter of the penalised (smoothing-spline) fit
$\int g''(t)^2 dt$	Roughness penalty: total curvature of the fit g
df_λ	Effective degrees of freedom, $\text{trace}(S_\lambda)$, where $\hat{g} = S_\lambda y$
s	Span: fraction of nearest neighbours used by LOESS
$f_j(X_j)$	Smooth component function of predictor X_j in a GAM

Sources and references

Primary textbook

James, G., Witten, D., Hastie, T., Tibshirani, R., & Taylor, J. (2023). *An Introduction to Statistical Learning, with Applications in Python*. Springer.

About these slides

Each method is introduced with motivation, intuition, formal definition, and worked examples. Slide titles are descriptive.

Learning objectives for this chapter

By the end of this chapter you will be able to:

- **Fit** polynomial and step-function regressions and discuss their respective failure modes.
- **Construct** regression splines, smoothing splines, and local regressions; **choose** smoothness by cross-validation.
- **Articulate** the boundary problem of cubic splines and the fix offered by natural splines.
- **Build** a generalised additive model (GAM) over multiple predictors and **interpret** its component functions.
- **Reason** about when nonlinear methods help and when they merely add variance.

How to use this deck — self-study and revision

Attempt each exercise *before* its solution; run `chapter_07_lab.ipynb` alongside the lab section. Revise from the back — key formulas, vocabulary, self-check (each question names its page) — then redo the exercises; the appendix is fully solved.

Roadmap of this chapter

A ladder of flexibility

- 1 Polynomials.
- 2 Step functions.
- 3 Regression splines (cubic, natural).
- 4 Smoothing splines.
- 5 Local regression (LOESS).
- 6 Generalised additive models (GAMs).

All of these are linear regression with cleverly chosen basis functions.

Outline

- 1 Why this chapter matters
- 2 Polynomials
- 3 Splines
- 4 Smoothing
- 5 GAMs
- 6 Python Lab
- 7 Summary

Why move beyond linearity?

Linear regression is interpretable and stable, but the world is rarely linear. Whenever the relationship between X and Y bends, a linear model leaves bias on the table.

The strategy of this chapter

Replace the strict linear form $\beta_j X_j$ with smooth functions $f_j(X_j)$, while keeping the additive structure that makes linear regression so easy to interpret.

Three reasons nonlinearity matters

Concrete examples

- Wage data peaks in middle age and then drops — not linear.
- Sales have diminishing returns to advertising spend — not linear.
- Survival probability is bounded in $(0, 1)$ — not linear.

Takeaway

Every method in this chapter is just linear regression with a cleverly chosen *basis expansion* of the original predictors. The OLS machinery still applies.

Live demo — §1 of the notebook

Switch to `chapter_07_lab.ipynb` §1 and run the two polynomial cells: a raw-power degree-4 OLS on Wage ($R^2 \approx 0.086$) whose grid plot shows exactly that middle-age wage hump.

Outline

- 1 Why this chapter matters
- 2 Polynomials**
- 3 Splines
- 4 Smoothing
- 5 GAMs
- 6 Python Lab
- 7 Summary

Polynomial regression: the simplest extension

Replace X with a basis of powers.

Model

$$Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \cdots + \beta_d X^d + \varepsilon.$$

How to read this

Still a *linear* model in the coefficients — all OLS machinery applies (standard errors, F -tests, CV). We are simply augmenting the design matrix.

Polynomial regression: the boundary problem

Beware of high degree

Degrees beyond 4 rarely improve fit but routinely produce wild oscillations near the boundaries of the data range. Splines are the safer alternative.


Why the wiggles?

A single global polynomial must use *one* formula everywhere. Each extra degree buys another turning point, and where data are sparse (the edges) those bends are barely constrained — so the curve swings and its confidence band balloons.

Takeaway

A useful tool for capturing mild curvature, especially with degree 2 or 3. Beyond that, switch to splines.

Wage data: degree-4 polynomial



Takeaway

Left: the degree-4 fit with its pointwise band; right: the fitted $\Pr(\text{wage} > 250 \mid \text{age})$. The curve captures the rise-then-plateau, but bends and widens at the sparse upper edge.

Step functions: piecewise constant

An alternative when the relationship is naturally piecewise constant.

Model


Bin X into $K + 1$ ranges and fit a constant in each:

$$Y = \beta_0 + \beta_1 1\{c_1 \leq X < c_2\} + \cdots + \beta_K 1\{c_K \leq X\} + \varepsilon.$$

Pros and cons

Easy to interpret (“income bracket”), but discontinuous at the breakpoints — often visually ugly. Useful when natural categories exist (age groups, salary tiers), or when banding is chosen deliberately for auditability — at the price of a visible jump for two cases either side of a cut-off.

Step function fit on Wage



Takeaway

Piecewise-constant: flat within each band, jumping at every cutpoint. The level shifts are captured, but the discontinuities are an artefact of binning.

Exercise 7.1 — Evaluating a fitted cubic [Math]

Exercise (≈ 5 min) — try it before the solution

Consider a cubic polynomial regression of wage on age using the raw basis $\{1, x, x^2, x^3\}$.

- 1 Write the design-matrix row for a single observation with $x = 2$.
- 2 The fitted model is $\hat{f}(x) = 50 + 4x - 0.5x^2 + 0.1x^3$. Predict the response at $x = 2$ and $x = 4$.
- 3 Degree three sounds nonlinear. Explain why this is still a *linear* model.

Solution — Exercise 7.1 (1/4)

Solution

(a) **Design-matrix row.** Polynomial regression is ordinary OLS on a design matrix whose columns are powers of x , so each observation contributes the row of its basis values. With basis $\{1, x, x^2, x^3\}$ and $x = 2$:

$$[1, 2, 2^2, 2^3] = [1, 2, 4, 8].$$

For n observations the full design matrix stacks such rows: $X_{i\cdot} = [1, x_i, x_i^2, x_i^3]$.

Solution — Exercise 7.1 (2/4)

Solution

(b) Predict by substituting term by term into $\hat{f}(x) = 50 + 4x - 0.5x^2 + 0.1x^3$:

$$\hat{f}(2) = 50 + 4(2) - 0.5(4) + 0.1(8) = 50 + 8 - 2 + 0.8 = 56.8,$$

$$\hat{f}(4) = 50 + 4(4) - 0.5(16) + 0.1(64) = 50 + 16 - 8 + 6.4 = 64.4.$$

Interpretation: between $x = 2$ and $x = 4$ the prediction rises by $64.4 - 56.8 = 7.6$; the curve still climbs here, but the negative x^2 term is already damping the growth.

Solution — Exercise 7.1 (3/4)

Solution

(c) **Why this is still a linear model.** “Linear” refers to the *coefficients*, not to the shape of the fitted curve. Writing

$$\hat{f}(x) = \beta_0 \cdot 1 + \beta_1 \cdot x + \beta_2 \cdot x^2 + \beta_3 \cdot x^3,$$

we see $\hat{f}$ is a linear combination of *known, pre-computed* regressors $1, x, x^2, x^3$. Hence all of the OLS machinery applies unchanged: closed-form $\hat{\beta}$, standard errors, t - and F -tests, cross-validation. The nonlinearity lives in x , not in the parameters.

Solution — Exercise 7.1 (4/4)

Solution

Interpretation: this is the basis-function idea behind the whole chapter — keep the linear-regression machinery, change only the columns of the design matrix.

Common mistake

Predicting $\hat{f}(4)$ by doubling $\hat{f}(2)$. Polynomial predictions do not scale linearly in x ; every power term must be substituted separately.

Exercise 7.2 — Step functions and dummy coding [Concept]

Exercise (≈ 5 min) — try it before the solution

We fit a piecewise-constant (step-function) regression of wage on age using cutpoints at 35 and 50, giving three bins: $(-\infty, 35)$, $[35, 50)$, $[50, \infty)$.

- 1 How many dummy variables do we create, and why is one bin omitted?
- 2 Write the fitted model with an intercept.
- 3 With $\hat{\beta}_0 = 40$, $\hat{\beta}_1 = 15$ (bin $[35, 50)$), $\hat{\beta}_2 = 25$ (bin $[50, \infty)$), predict wage for ages 30, 40, and 60.
- 4 State one limitation of this approach.

Live demo — §2 of the notebook

Switch to `chapter_07_lab.ipynb` §2 and run the binning cell: four age bins give mean wages $94.2 \rightarrow 118.2 \rightarrow 117.8 \rightarrow 101.8$ — the hump no straight line can produce, at the price of a jump at age 33.5.

Solution — Exercise 7.2 (1/3)

Solution

(a) Two dummies for three bins. Define $C_1(x) = 1[35 \leq x < 50]$ and $C_2(x) = 1[x \geq 50]$. The first bin $(-\infty, 35)$ is the *reference category*, absorbed into the intercept. *Why omit one bin?* With an intercept, the column of ones equals the sum of all three bin dummies, so including all three would make the design matrix perfectly collinear — the “dummy-variable trap”. In general: m bins $\Rightarrow m - 1$ dummies plus an intercept.

(b) Fitted model.

$$y = \beta_0 + \beta_1 C_1(x) + \beta_2 C_2(x) + \varepsilon.$$

Reading the coefficients: β_0 is the mean wage in the reference bin; β_1 and β_2 are the mean differences of bins 2 and 3 relative to that reference — not relative to each other.

Solution — Exercise 7.2 (2/3)

Solution

(c) **Predictions.** Locate each age's bin, switch on its dummy, and add:

- Age $30 < 35$: reference bin, $C_1 = C_2 = 0 \Rightarrow \hat{y} = 40$.
- Age $40 \in [35, 50)$: $C_1 = 1, C_2 = 0 \Rightarrow \hat{y} = 40 + 15 = 55$.
- Age $60 \geq 50$: $C_1 = 0, C_2 = 1 \Rightarrow \hat{y} = 40 + 25 = 65$.

Interpretation: predicted wage steps from 40 to 55 at age 35 and to 65 at age 50 — everyone inside a bin gets the same prediction.

Solution — Exercise 7.2 (3/3)

Solution

(d) Limitation. The fit is constant within each bin and jumps at the cutpoints: it ignores any trend inside a bin, and the cutpoints are chosen arbitrarily. Unless the true relationship really has breakpoints, this misses smooth structure.

Common mistake

Reading $\beta_2 = 25$ as the jump from bin 2 to bin 3. Both coefficients compare to the *reference* bin; the bin-2-to-bin-3 step is $\beta_2 - \beta_1 = 10$.

Outline

- 1 Why this chapter matters
- 2 Polynomials
- 3 Splines**
- 4 Smoothing
- 5 GAMs
- 6 Python Lab
- 7 Summary

Basis functions: a unifying view

Polynomials and step functions are special cases of:

Basis-function regression

$$Y = \beta_0 + \sum_{j=1}^K \beta_j b_j(X) + \varepsilon,$$

for fixed basis functions $b_1, \dots, b_K$.

How to read this

Read it as: each $b_j(\cdot)$ is a *fixed* transformation chosen in advance, and β_j is its ordinary regression weight — so once the b_j are set, this is plain OLS. Special cases:

- Polynomials: $b_j(X) = X^j$.
- Step functions: $b_j(X) = 1\{c_j \leq X < c_{j+1}\}$.
- Splines: piecewise polynomials joined smoothly at chosen knots.

Regression splines: idea

Splines combine the flexibility of polynomials with the locality of step functions.

Construction

Split the range of X at K *knots* $\xi_1 < \dots < \xi_K$. Fit a polynomial of degree d on each piece, joined at the knots so the function and its first $d - 1$ derivatives are continuous.



Practical defaults

Cubic splines ($d = 3$) are the standard choice — they are smooth visually. Total degrees of freedom: $K + d + 1$.

Live demo — §3 of the notebook

Switch to `chapter_07_lab.ipynb` §3 and run the `SplineTransformer` cell: a cubic B-spline basis with 6 knots, fitted by plain least squares — the curve bends like the polynomial but stays smooth with no jumps.


Why the continuity constraints? Broken vs. smooth joins



Takeaway

A regression spline is *not* one global polynomial: it is separate low-degree pieces. All the work is in the join — forcing equal value, slope and curvature at each knot makes the seam invisible while each piece stays flexible.

How a cubic spline is built: continuity at the knots



Takeaway

Each region is a free cubic, yet at every knot the pieces share value, slope and curvature — so the joins are invisible.

The natural-spline boundary fix

Boundary problem

Cubic splines flare beyond extreme knots. The CIs at the boundaries become unreliably wide.

Why the edges misbehave

Past the outermost knots the fit is still a free cubic, but almost no data pin it down there. A cubic can curve sharply, so a small change in a boundary point swings the whole tail — inflating variance exactly where evidence is thinnest.

Natural cubic splines

Constrain the spline to be *linear* past the boundary knots. The resulting fits are stable at the edges with negligible bias inside.

Choosing knots and degrees of freedom

Two practical decisions

- **How many knots?** Choose by cross-validation. Modest df (4–8) usually suffices.
- **Where?** Equal quantiles of X usually work well. With small n , place knots manually at obvious change-points.


Why cubic and not quintic?

Cubic is the lowest degree that gives visually smooth fits (matches second derivative). Higher degrees rarely improve the fit but add parameters and instability.

Takeaway

The knot count is the real tuning parameter: more knots track every wiggle in the data, fewer give a smooth curve that misses some of it. Choose it by cross-validation, not by eye.

Cubic vs. natural spline on Wage



Takeaway

Same data, two fits: the cubic spline's error bands blow up at the edges while the natural spline's stay controlled — boundary behaviour is the deciding difference.

How do we actually fit one? The truncated-power basis

A cubic spline is just OLS on the right columns. Start from a cubic, then add *one* extra term per knot:

The basis, with K knots $\xi_1 < \dots < \xi_K$


$$1, x, x^2, x^3, (x - \xi_1)_+^3, \dots, (x - \xi_K)_+^3, \quad (x - \xi)_+^3 = \begin{cases} (x - \xi)^3 & x > \xi, \\ 0 & \text{otherwise.} \end{cases}$$

Why this works, and what it costs

Each truncated term is *zero left of its knot*, so it switches on exactly where its piece begins — the locality a global polynomial can never have. It changes only the *third* derivative at ξ , so value, slope and curvature stay continuous automatically. Count the columns: $4 + K$ coefficients, i.e. $K + 4$ degrees of freedom, matching the count you are about to derive.

Source: Section 7.4.3, equations (7.9)–(7.10) — ISLP, James et al. (2023).

The truncated-power basis, drawn



Takeaway

The same basis, plotted: four global polynomial terms plus one truncated cubic per knot ($K = 3$ knots at the Wage age quartiles) $= K + 4 = 7$ basis functions. Note where each orange curve leaves the axis — that is its knot.

Exercise 7.3 — Degrees of freedom of a cubic spline [Math]

Exercise (≈ 5 min) — try it before the solution

Consider a cubic regression spline with K interior knots, fitted so that the function and its first two derivatives are continuous at each knot.

- 1 If we fit a *separate* cubic in each of the $K + 1$ regions, how many free parameters is that?
- 2 How many constraints do the continuity conditions impose in total?
- 3 Show that the net degrees of freedom equal $K + 4$.
- 4 How many basis functions does the truncated-power basis use (including the intercept)? Confirm consistency.

Live demo — §6 of the notebook

Switch to `chapter_07_lab.ipynb` §6 and run the basis-building cell: it constructs the truncated-power basis on Wage ($K = 3$ knots at the age quartiles) and confirms the design matrix has rank $7 = K + 4$ — the count you are about to derive.

Solution — Exercise 7.3 (1/3)

Solution

(a) **Free parameters.** K knots split the line into $K + 1$ regions. Each region carries its own cubic $a_j + b_jx + c_jx^2 + d_jx^3$, i.e. 4 coefficients, so fitting them separately costs

$$4 \times (K + 1) = 4K + 4 \text{ parameters.}$$

(b) **Constraints.** At each interior knot we require the two adjacent pieces to agree in *value*, *slope*, and *curvature* (g, g', g'' continuous): 3 constraints per knot, hence

$$3 \times K = 3K \text{ constraints in total.}$$

Why stop at the second derivative? Matching g''' as well would force adjacent cubics to be identical — the “spline” would collapse to one global cubic. Continuity up to g'' is exactly what the eye perceives as smooth.

Solution — Exercise 7.3 (2/3)

Solution

(c) **Net degrees of freedom.** Each independent linear constraint removes one free parameter:

$$\underbrace{(4K + 4)}_{\text{free params}} - \underbrace{3K}_{\text{constraints}} = K + 4.$$

Sanity check: $K = 0$ gives 4 — a single cubic, as it must. Each additional knot buys exactly *one* extra degree of freedom.

Solution — Exercise 7.3 (3/3)

Solution

(d) **Basis count.** The truncated-power basis is

$$1, x, x^2, x^3, (x - \xi_1)_+^3, \dots, (x - \xi_K)_+^3,$$

which is $4 + K$ functions — one regression coefficient each, exactly matching the $K + 4$ df from (c). Each truncated term adds one knot while preserving continuity of g, g', g'' , so the two counting arguments agree.

Common mistake: counting 4 constraints per knot (including g''') or forgetting the intercept among the basis functions — both break the $K + 4$ tally.

Exercise 7.4 — Natural vs. cubic splines [Concept]

Exercise (≈ 5 min) — try it before the solution

- 1 What extra constraints turn a cubic spline into a *natural* cubic spline?
- 2 A cubic spline with K knots has $K + 4$ df. How many df does a natural cubic spline with K knots have?
- 3 Why do natural splines behave better near the boundaries of the data?
- 4 What does this imply for the width of pointwise confidence bands at the extremes of x ?

Solution — Exercise 7.4 (1/3)

Solution

(a) Extra constraints. A natural spline is required to be *linear* beyond the two boundary knots. A function is linear exactly when its second and third derivatives vanish, so the requirement translates into

$$g''(\xi) = 0 \quad \text{and} \quad g'''(\xi) = 0 \quad \text{at and beyond each boundary knot,}$$

i.e. 2 extra constraints at each of the 2 boundaries: 4 additional constraints in total.

(b) Degrees of freedom. Start from the cubic-spline count and subtract one df per constraint:

$$(K + 4) - 4 = K.$$

Interpretation: a natural spline with K knots is exactly as complex as a linear model with K coefficients — its flexibility budget is spent entirely in the interior, where the data live.

Solution — Exercise 7.4 (2/3)

Solution

(c) **Why natural splines behave at the boundaries.** Past the outer knots an unrestricted cubic spline is still a free cubic, but almost no data pin it down there: a handful of extreme observations steer four polynomial coefficients, so small data changes swing the whole tail — high variance exactly where evidence is thinnest. Forcing linear tails removes those unsupported cubic wiggles.

Solution — Exercise 7.4 (3/3)

Solution

(d) Consequence for confidence bands. The pointwise band width tracks the variance of $\hat{g}(x)$. Because fewer parameters are spent at the extremes and the tail shape is constrained to be linear, natural-spline bands are substantially *narrower* at the boundaries than cubic-spline bands, at the price of a negligible bias increase in the interior.

Common mistake

Thinking the natural spline is linear *between* the boundary knots too — inside the boundary knots it is every bit as flexible a cubic spline; only the two tails are straightened.

Outline

- 1 Why this chapter matters
- 2 Polynomials
- 3 Splines
- 4 Smoothing**
- 5 GAMs
- 6 Python Lab
- 7 Summary

Smoothing splines: let the data choose

Instead of choosing knots manually, place a knot at *every* data point and let a penalty control smoothness.

Penalised RSS

$$\min_g \sum_{i=1}^n (y_i - g(x_i))^2 + \lambda \int g''(t)^2 dt.$$

Takeaway

Two terms pull opposite ways: the RSS wants a curve through every point, the penalty wants a straight line. λ sets the exchange rate — and over-smoothing flattens exactly the peaks that matter most.

Reading the smoothing-spline objective


How to read this

Term	Role
$\sum (y_i - g(x_i))^2$	Training fit
$\int g''(t)^2 dt$	Penalty for “wiggleness” (curvature)
λ	Smoothness tuning parameter

Three observations

- The minimiser is a natural cubic spline with knots at every x_i .
- $\lambda = 0$: interpolation. $\lambda \rightarrow \infty$: linear fit.
- The fit is linear in y — $\hat{g} = S_\lambda y$ — so smoothness is reported as the **effective degrees of freedom** $\text{df}_\lambda = \text{trace}(S_\lambda)$, running from n (interpolation) down to 2 (a straight line). Quote df_λ , not λ : it is comparable across data sets.


Smoothing spline on Wage



Takeaway

Two ways to fix smoothness: a preset 16 df vs. the ≈ 6.8 effective df chosen by CV. More effective df means a wigglier fit; CV finds the sweet spot.

Smoothing spline flexibility: λ is the dial



Takeaway

Same Wage data, two penalties: $\lambda = 10^4$ gives a calm trend (≈ 3.7 effective df) while $\lambda = 0.1$ (≈ 43 df) chases noise — λ is the flexibility dial, tuned by CV/LOOCV.

Local regression (LOESS): a different philosophy

Instead of one global function, fit a small model around each point.


Recipe at each query point x_0

- 1 Find the fraction s (“span”) of points closest to x_0 .
- 2 Weight them by a kernel that decays with distance.
- 3 Fit a low-degree polynomial by weighted least squares.
- 4 Evaluate at x_0 ; repeat over the grid of x_0 .

How to read this

Span s controls smoothness — the analogue of λ . LOESS is excellent for one or two predictors and dies fast as p grows.

LOESS on Wage



Takeaway

Each curve uses a different span: 0.2 tracks local wiggles, 0.7 is smoother. Span plays the role of λ , trading bias against variance.

Mini-check: spline choices

Quick self-assessment

- 1 Why prefer a natural cubic spline over a regular cubic spline at the boundaries?
- 2 How is a smoothing spline's λ related to its effective degrees of freedom?
- 3 In LOESS, what happens to bias as the span shrinks?

Answers

1. Natural splines are constrained linear past the boundary knots, avoiding wild boundary oscillation. 2. Larger $\lambda \Rightarrow$ smaller df (more shrinkage, smoother fit). 3. Bias decreases (the local fit is more local), but variance increases.

Exercise 7.5 — Smoothing splines and λ [Concept]

Exercise (≈ 5 min) — try it before the solution

A smoothing spline chooses g to minimise

$$\sum_{i=1}^n (y_i - g(x_i))^2 + \lambda \int g''(t)^2 dt.$$

- 1 Interpret the penalty term. What does it measure?
- 2 Describe the fit as $\lambda \rightarrow 0$ and as $\lambda \rightarrow \infty$.
- 3 Define the effective degrees of freedom df_λ and give its range.
- 4 How is λ chosen in practice?

Solution — Exercise 7.5 (1/3)

Solution

(a) The penalty measures curvature. $g''(t)$ is the rate of change of the slope at t ; squaring removes the sign and integrating adds curvature up over the whole range, so $\int g''(t)^2 dt$ is the total *roughness* of g . A straight line has $g'' \equiv 0$ and pays no penalty; a wiggly curve pays a lot. The penalty therefore discourages wiggleness.

(b) The two extremes. As $\lambda \rightarrow 0$ curvature becomes free, so the minimiser drives the RSS term to 0: g *interpolates* the data — extremely wiggly, low bias, high variance. As $\lambda \rightarrow \infty$ any curvature is infinitely expensive, forcing $g'' = 0$ everywhere, so g collapses to the ordinary least-squares *line*.

Interpretation: λ sweeps a continuous bridge between interpolation and linear regression — it is the bias–variance dial of the smoothing spline.

Solution — Exercise 7.5 (2/3)

Solution

(c) Effective degrees of freedom. For fixed λ the fit is linear in y : $\hat{g} = S_\lambda y$ for a smoother matrix S_λ . In OLS the analogous “hat” matrix has trace equal to the number of parameters, which motivates the definition

$$\text{df}_\lambda = \text{trace}(S_\lambda),$$

ranging from n (at $\lambda = 0$, interpolation) down to 2 (at $\lambda \rightarrow \infty$, slope + intercept of the linear fit).

Solution — Exercise 7.5 (3/3)

Solution

(d) Choosing λ . By cross-validation. LOOCV is essentially free because each leave-one-out residual follows from the *single* full fit via the leverage $\{S_\lambda\}_{ii}$:

$$\text{RSS}_{\text{cv}}(\lambda) = \sum_{i=1}^n \left(\frac{y_i - \hat{g}_\lambda(x_i)}{1 - \{S_\lambda\}_{ii}} \right)^2.$$

Common mistake

Picking λ by minimising *training* RSS — that always selects $\lambda = 0$ (interpolation). Only out-of-sample criteria can trade fit against roughness.

Outline

- 1 Why this chapter matters
- 2 Polynomials
- 3 Splines
- 4 Smoothing
- 5 GAMs**
- 6 Python Lab
- 7 Summary

From one predictor to many: GAMs

Everything we have done so far is for one predictor at a time. Generalised additive models lift the techniques into the multivariate setting.

The GAM model

$$Y = \beta_0 + \sum_{j=1}^p f_j(X_j) + \varepsilon,$$

where each f_j is a smooth function (spline or LOESS).

Symbol by symbol

β_0 is the overall level; each f_j is a smooth, data-driven curve for predictor X_j ; summing them keeps the model additive. Setting $f_j(X_j) = \beta_j X_j$ recovers ordinary linear regression — the GAM simply lets each term bend.

Why GAMs are so useful


Three reasons

- 1 Generalises linear regression: $f_j(X_j) = \beta_j X_j$ is the linear special case.
- 2 Avoids the curse of dimensionality via the additive structure.
- 3 Each f_j can be plotted separately — excellent interpretability.

Limitation

By construction, GAMs miss interactions. Add tensor-product terms $f_{jk}(X_j, X_k)$ explicitly, or move to trees (Ch. 8).

GAM fit on Wage




Each panel: estimated f_j for a different predictor, with pointwise confidence bands.

Source: Figure 7.12 — ISLP, James et al. (2023).

GAM component functions fitted on Wage

GAM component functions: $\text{wage} \sim s(\text{age}) + s(\text{year}) + \text{education}$



Takeaway

Reading each partial effect: wage rises with age then falls after ~ 60 ; it drifts up gently with year; and it increases monotonically with education. The additive structure lets us plot and interpret one predictor at a time.

GAMs for classification

Same idea inside a GLM:

Logistic GAM

$$\log \frac{p(X)}{1 - p(X)} = \beta_0 + \sum_{j=1}^p f_j(X_j).$$

Reading the model

The left side is the *log-odds* of the positive class; each smooth f_j shifts those log-odds up or down as X_j changes. It is an additive logistic model — interpret every f_j on its own, just as in a linear GLM.

Combines logistic regression's interpretability with smooth per-predictor effects.

Implementations: `pyGAM`, `statsmodels.gam`.

Extended Exercise 7.2 — A GAM for wage [Integrative]

Extended exercise (≈ 15 min) — try it before the solution

Fit a generalised additive model

$$\text{wage} \approx \beta_0 + s(\text{age}) + \text{year} + f(\text{education}),$$

with a smooth spline in age, a (near-)linear term in year, and education as a factor.

- 1 Fit it (e.g. pyGAM) and state how a GAM is estimated.
- 2 Interpret each partial effect.
- 3 Which term is genuinely non-linear? Discuss the effective degrees of freedom (EDF) of each term.
- 4 State the advantages of the additive form and its key limitation. Give an interaction it would miss and how to add it.

Solution — Extended Exercise 7.2 (1/4)

Solution

(a) Fit & algorithm. Encode education to integer codes, then fit a spline in age, a linear term in year and a factor in education:

```
import pandas as pd
from pygam import LinearGAM, s, l, f

W = pd.read_csv('../ALL CSV FILES - 2nd Edition/Wage.csv') # load Wage
W['edu'] = W['education'].astype('category').cat.codes # levels -> 0..4
X = W[['age', 'year', 'edu']].values # cols 0, 1, 2 in that order
y = W['wage'].values # response
gam = LinearGAM(s(0) + l(1) + f(2)).fit(X, y) # smooth, linear, factor
gam.summary() # EDF per term: age ~15, year ~1, edu 4
```

A GAM is fit by *backfitting*: cycle through the terms, repeatedly re-estimating each f_j on the partial residual $y - \beta_0 - \sum_{k \neq j} f_k(X_k)$ until the fit stabilises — each smooth is refitted to whatever the other terms leave unexplained.

Solution — Extended Exercise 7.2 (2/4)

Solution

(b) **Partial effects** (each holds the others fixed):

- age: a pronounced *hump* — wage climbs steeply to about age 40, plateaus through the 50s, then declines after ~ 60 .
- year: a gentle, near-linear *increase* — slow wage growth over the sample window.
- education: a monotone *staircase* — each higher band pays more, “Advanced Degree” far above “<HS Grad”.

Solution — Extended Exercise 7.2 (3/4)

Solution

(c) Non-linearity & EDF. Only age is genuinely non-linear: its smooth spends many effective degrees of freedom (≈ 15 at pyGAM's default penalty, dropping to ≈ 10 once λ is tuned by `gam.gridsearch`). year enters with $\text{EDF} \approx 1$ (essentially a straight line); education as a factor contributes $\text{df} = (\text{levels} - 1) = 4$. The EDF is how a GAM quantifies “how wiggly” each term is.

Common mistake

Reading a partial-effect curve as an absolute wage level. Each panel is centred around zero; only its *shape* and within-panel differences are meaningful, with all other predictors held fixed.

Solution — Extended Exercise 7.2 (4/4)

Solution

(d) Advantages of additivity.

- Each term is estimated and *plotted separately*, so effects are directly interpretable.
- The additive structure sidesteps the curse of dimensionality (no p -dimensional smoother).
- It strictly generalises linear regression: a linear term is the special case $f_j(X_j) = \beta_j X_j$.

Key limitation — no interactions. By construction the shape of $s(\text{age})$ is the *same* at every education level. If the wage–age profile were steeper for college graduates than for high-school graduates (an $\text{age} \times \text{education}$ interaction), the additive GAM could not represent it. *Fix:* add a bivariate/tensor term $f(\text{age}, \text{education})$, fit separate age-smooths per group, or move to trees / boosting (Ch. 8), which capture interactions automatically.

Outline

- 1 Why this chapter matters
- 2 Polynomials
- 3 Splines
- 4 Smoothing
- 5 GAMs
- 6 Python Lab**
- 7 Summary

Python lab — run it live in the notebook

Companion notebook: `chapter_07_lab.ipynb`

The code in this section is meant to be **demonstrated live**. Open the companion Jupyter notebook and run it cell by cell, section by section:

- §1, *Polynomial regression*, and §2, *Step functions*;
- §3, *Regression splines*, and §4, *LOESS*;
- §5, *GAM*, and §6, *Where the degrees of freedom come from*.

Data loads via the ISLP package or the bundled CSV files; §7, *Exercises*, closes the notebook.

Polynomial regression

```
import statsmodels.api as sm
from ISLP import load_data
from ISLP.models import poly, ModelSpec as MS

Wage = load_data("Wage") # 3000 workers: wage, age, year
design = MS([poly("age", degree=4)]).fit_transform(Wage) # 4 age cols
res = sm.OLS(Wage["wage"], design).fit() # plain OLS on that design
print(res.summary()) # one coefficient per basis column
```

What the design matrix actually holds

`fit_transform` is what builds the columns OLS sees: an intercept plus *four* columns for age. `poly` returns *orthogonal* polynomials — rotated combinations of age , age^2 , age^3 , age^4 that span the same space, so the fitted curve and $R^2 \approx 0.086$ match raw powers exactly, but the columns are uncorrelated. Judge the term by the fitted curve, never by one column's t -statistic.

Regression splines

```
from ISLP.models import bs, ns
```

```
design = MS([bs("age", df=6)]).fit_transform(Wage) # 6 cubic cols  
design2 = MS([ns("age", df=4)]).fit_transform(Wage) # 4 cols, linear tails  
res = sm.OLS(Wage["wage"], design2).fit() # still ordinary least squares  
print(res.summary()) # one weight per basis function
```

df and knots are two ways to say the same thing

Each call returns a matrix whose columns are the basis functions $b_j(\text{age})$ of the slide before. `df` says *how many columns*; the knots are then placed for you at quantiles of age. Naming the knots instead gives the identical fit, because cubic-spline $\text{df} = K + 4$ and natural-spline $\text{df} = K$ — which is why `ns` at the same `df`, spending nothing on the tails, can afford more knots where the data actually are.

Local regression with statsmodels

```
import statsmodels.nonparametric.api as nps

l = nps.lowess(endog=Wage["wage"], exog=Wage["age"], # y first, then x
               frac=0.2, return_sorted=True) # span: 20% of the points

import matplotlib.pyplot as plt
plt.scatter(Wage["age"], Wage["wage"], s=4, alpha=0.3) # raw data
plt.plot(l[:, 0], l[:, 1], color="orange") # col 0 = age, col 1 = fit
plt.show()
```

No coefficients come back — just the curve

One $(x, \hat{y})$ row per point, sorted by x so it can be plotted.

Live demo — §4 of the notebook

Switch to `chapter_07_lab.ipynb` §4 and run the `lowess` cell: with span 0.2 each local fit uses 20% of the 3000 workers, and the curve traces the same rise–plateau–fall as the splines.

A simple GAM with pyGAM

```
from pygam import LinearGAM, s
```

```
X = Wage[["year", "age"]].values # col 0 = year, col 1 = age
```

```
y = Wage["wage"].values # response
```

```
gam = LinearGAM(s(0) + s(1)).fit(X, y) # s(j): smooth of column j
```

```
gam.summary() # EDF per term: year ~7, age ~14
```

Reading the EDF column

`s(0)` and `s(1)` index *columns* of `X`. A term's effective df says how much flexibility the penalty let it use: 1 is a straight line, year ≈ 7 a gentle drift, age ≈ 14 the pronounced hump.

Live demo — §5 of the notebook

Switch to `chapter_07_lab.ipynb` §5 and run the two GAM cells: the summary reports ≈ 21.2 total effective df (age's smooth alone spends ≈ 14), and the partial-dependence panels show the age hump next to the gentle year drift.

Exercise 7.6 — Natural spline / GAM on Wage [Python]

Exercise (≈ 5 min) — try it before the solution

Using `Wage.csv`:

- 1 Fit wage on a natural cubic spline of age with 5 df via OLS.
- 2 Fit a GAM $\text{wage} \sim s(\text{age}) + s(\text{year})$.
- 3 Describe the estimated effect of age.

Run this live

The worked solution sits in `chapter_07_lab.ipynb` under *Lecture exercises — worked Python solutions*: the cell headed *Exercise 7.6 — Natural spline / GAM on Wage [Python]*.

Solution — Exercise 7.6 (1/3)

Solution

(a) Natural cubic spline via patsy's `cr()` basis — `cr()` expands age into 5 natural-spline columns (linear tails built in), so the fit is still ordinary OLS:

```
import pandas as pd
import statsmodels.formula.api as smf

Wage = pd.read_csv('../ALL CSV FILES - 2nd Edition/Wage.csv') # load
fit = smf.ols('wage ~cr(age, df=5)', data=Wage).fit() # 5 ns columns
print(fit.summary()) # read the joint F, not single t's
```

What the output shows: 6 coefficients (intercept + 5 basis columns), $R^2 \approx 0.087$, and an overall F -test with $p \approx 6 \times 10^{-58}$ — the spline terms are jointly highly significant, even though age alone explains only a modest share of wage variance. `df=5` is the *entire* complexity choice: it fixes how many basis columns are built, and the knots then go at quantiles of age.

Solution — Exercise 7.6 (2/3)

Solution

(b) A GAM with smooth terms via pyGAM — each $s()$ is a penalised spline whose wiggleness is measured in effective degrees of freedom (EDF):

```
import numpy as np
from pygam import LinearGAM, s

X = Wage[['age', 'year']].values # col 0 = age, col 1 = year
y = Wage['wage'].values # response
gam = LinearGAM(s(0) + s(1)).fit(X, y) # a penalised spline per column
gam.summary() # EDF per term: how wiggly each one got
```

What the output shows: with pyGAM's default penalty the age smooth uses ≈ 15 EDF and year ≈ 6 — both flexible; tuning λ by `gam.gridsearch(X, y)` shrinks the smooths toward simpler shapes. Note the contrast with (a): there *you* fixed $df = 5$; here λ decides, and the EDF merely *reports* the flexibility it granted.

Solution — Exercise 7.6 (3/3)

Solution

(c) *Interpretation.* The fitted effect of age is clearly nonlinear: wage rises steeply from the late teens to about age 40, is roughly flat (a plateau) between 40 and 60, then declines gently after 60. A straight line would miss this rise-then-plateau-then-fall shape, which is why a spline / GAM is preferable. The year effect is mildly increasing (slow wage growth over time).

How to read the fit: plot `gam.partial_dependence()` per term rather than staring at coefficients — the spline basis coefficients have no individual meaning; only the fitted curve does.

Common mistake

Interpreting single spline-basis coefficients (or their *t*-tests) as “the effect of age”. Judge the term by its fitted curve and a *joint* test of all its basis columns.

Extended Exercise 7.3 — Three fits for wage vs. age [Python]

Extended exercise (≈ 15 min) — try it before the solution

Using `Wage.csv`, model wage as a function of age three ways and compare:

- 1 a **polynomial**, choosing the degree by nested ANOVA F -tests (or 5-fold CV);
- 2 a **step function** (bin age into several ranges);
- 3 a **natural cubic spline** with about 5 df.

Overlay the three fitted curves on a scatterplot, report the chosen complexity, and interpret where the fits agree and where they differ.

Run this live

The worked solution sits in `chapter_07_lab.ipynb` under *Lecture exercises — worked Python solutions*: the cell headed *Extended Exercise 7.3 — Three fits for wage vs. age [Python]* — the nested ANOVA, the three fits and the overlay plot, ready to run.

Solution — Extended Exercise 7.3 (1/5)

Solution

Choose the polynomial degree by nested ANOVA.

```
import numpy as np, pandas as pd
import statsmodels.formula.api as smf
from statsmodels.stats.anova import anova_lm

W = pd.read_csv('../ALL CSV FILES - 2nd Edition/Wage.csv') # load Wage
terms = lambda d: '+' .join(f'np.power(age,{k})' for k in range(1, d+1)) # age..age^d
models = [smf.ols('wage ~'+terms(d), data=W).fit() for d in range(1, 6)] # nested: +1 term each
print(anova_lm(*models)[['df_diff', 'F', 'Pr(>F)']].round(4)) # row d: is degree d+1 worth it?
```

Each row of `anova_lm` tests degree d against $d+1$ — valid because the models are *nested*. The F -tests are highly significant up to degree 3 ($2 \rightarrow 3$: $p \approx 0.002$), *borderline* for $3 \rightarrow 4$ ($p \approx 0.05$) and clearly insignificant for $4 \rightarrow 5$ ($p \approx 0.37$). 5-fold CV of test MSE is essentially flat over degrees 3–4. **Chosen degree: 3 (cubic).**

Solution — Extended Exercise 7.3 (2/5)

Solution

Step function, natural spline, and the overlay.

```
import matplotlib.pyplot as plt
poly = smf.ols('wage ~'+terms(3), data=W).fit() # the degree chosen above
W['bin'] = pd.cut(W['age'], [17,25,35,45,55,65,81]) # 6 bins -> 5 dummies
step = smf.ols('wage ~C(bin)', data=W).fit() # a constant inside each bin
ns = smf.ols('wage ~cr(age, df=5)', data=W).fit() # 5 cols, linear tails

g = pd.DataFrame({'age': np.arange(18, 81)}) # one row per age to plot
g['bin'] = pd.cut(g['age'], [17,25,35,45,55,65,81]) # SAME cutpoints
plt.scatter(W['age'], W['wage'], s=4, alpha=.15, color='grey') # data
for m, lab in [(poly, 'poly (d=3)'), (step, 'step'), (ns, 'natural spline')]:
    plt.plot(g['age'], m.predict(g), label=lab) # three curves, one grid
plt.legend(); plt.xlabel('age'); plt.ylabel('wage'); plt.show()
```

Solution — Extended Exercise 7.3 (3/5)

Solution

Why the grid carries both columns: `predict` re-applies each formula, so `g` must contain `age` (for `poly, ns`) *and* the matching bin labels (for `step`), built with the *same* cutpoints as in training.

Solution — Extended Exercise 7.3 (4/5)

Solution

Reading the overlay.

- *Where they agree:* in the data-dense interior (ages ~ 25 – 70) all three fits trace the same rise–plateau–fall — wage climbs steeply to about age 40, is roughly flat through the 50s, then eases down.

Solution — Extended Exercise 7.3 (5/5)

Solution

- *Where they differ:* at the sparse boundaries. The polynomial can bend sharply at the extremes; the step function is blocky and *discontinuous* at the cut-points; the natural cubic spline stays smooth and, being linear beyond its outer knots, is the most stable at the edges.

Chosen complexity: a cubic polynomial (degree 3–4), a 6-bin step function, and a natural spline with ~ 5 df — all three explain a similar share of variance ($R^2 \approx 0.08$ – 0.09), but the natural spline is preferred for its boundary stability.

Common mistake

Picking the degree that maximises training R^2 — it rises mechanically with every added term; only the nested F -tests or CV can say when the gain is real.

Outline

- 1 Why this chapter matters
- 2 Polynomials
- 3 Splines
- 4 Smoothing
- 5 GAMs
- 6 Python Lab
- 7 Summary**

Chapter 7 in one slide

What we accomplished

A toolbox for adding flexibility to a regression while keeping the additive, interpretable structure of a linear model.

- Polynomials, step functions, regression splines, smoothing splines, LOESS.
- GAMs as the natural multivariate extension.
- Cross-validation chooses the smoothness in every method.

Key formulas at a glance

Polynomial $y_i = \beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \cdots + \beta_d x_i^d + \varepsilon_i$.

Step function $y_i = \beta_0 + \sum_{k=1}^K \beta_k C_k(x_i) + \varepsilon_i$, with bins C_k .

Cubic spline truncated-power basis $1, x, x^2, x^3, (x - \xi_k)_+^3$; K knots $\Rightarrow$ $df = K + 4$.

Natural spline linear beyond the boundary knots; $df = K$.

Smoothing spline $\min_g \sum_i (y_i - g(x_i))^2 + \lambda \int g''(t)^2 dt$.

Effective df $df_\lambda = \text{trace}(S_\lambda)$, where $\hat{g} = S_\lambda y$; larger $\lambda \Rightarrow$ *smaller* df_λ : from n (at $\lambda = 0$, interpolation) down to 2 (as $\lambda \rightarrow \infty$, a straight line).

GAM $y_i = \beta_0 + f_1(x_{i1}) + f_2(x_{i2}) + \cdots + f_p(x_{ip}) + \varepsilon_i$.

Backfitting re-estimate each f_j in turn on the partial residual $y - \beta_0 - \sum_{k \neq j} f_k(X_k)$, cycling through the terms until convergence.

Vocabulary checklist

Term	Meaning
Basis function	Fixed transformation $b_j(X)$ entering linearly
Knot	Breakpoint where pieces of a spline meet
Cubic / natural spline	Degree-3 piecewise polynomial / linear at boundaries
Smoothing spline	Penalised second-derivative fit
Effective degrees of freedom	Trace of the smoother matrix
LOESS / local regression	Weighted local polynomial
Span	Fraction of nearest neighbours used in LOESS
GAM	$Y = \beta_0 + \sum_j f_j(X_j) + \varepsilon$
Backfitting	Iterative algorithm for fitting GAMs

Methods comparison

Method	Strength	Weakness
Polynomial	Easy, captures curvature	Boundary oscillation
Step function	Interpretable	Discontinuous
Cubic spline	Smooth, controllable df	Boundary blow-up
Natural spline	Stable boundaries	Tiny extra bias inside
Smoothing spline	Auto-fitting smoothness	Black-box df
LOESS	Locally adaptive	Curse of dimensionality
GAM	Additive, interpretable	Misses interactions unless added

Decision rules of thumb

- For a single X and modest curvature: polynomial degree 2–4.
- For a single X with multiple regimes: regression splines with a handful of knots placed at quantiles.
- For boundary stability: *natural* cubic splines.
- For multiple predictors with little interaction: a GAM.
- Choose smoothness (df, λ , span) by cross-validation.
- For visible interactions: add tensor-product terms or move to trees (Ch. 8).

Common pitfalls

Don't

- 1 Use polynomials of degree > 4 on real data.
- 2 Place knots arbitrarily without checking residuals.
- 3 Forget that LOESS fails fast as p grows.
- 4 Compare GAM df with linear-regression df one-for-one.
- 5 Trust pointwise CIs of smoothing splines for inference.
- 6 Ignore extrapolation — splines outside the data range are unreliable.

Connections to the rest of the course

- Ch. 3 is the special case where every f_j is linear.
- Ch. 5 chooses $df / \lambda / \text{span by CV}$.
- Ch. 6 ridge / lasso can be applied to spline basis expansions.
- Ch. 8 trees are the nonparametric alternative when interactions dominate.
- Module A9's deep nets generalise GAMs by learning the basis.

Self-check questions

Each question names the slide that answers it — a wrong answer tells you where to read.

- 1 Why is a cubic spline preferred over a quintic spline? (p. 32)
- 2 Argue intuitively why smoothing splines have a unique minimiser. (p. 45)
- 3 How does the smoothing-spline penalty λ relate to df? (p. 46)
- 4 Sketch a GAM with two predictors. Where would you add an interaction? (pp. 57–58)
- 5 For LOESS, what happens to bias and variance as the span shrinks? (p. 49)
- 6 In the Wage data, would you expect a linear or a spline term in age to win? Why? (p. 60)

Seven things to remember

- 1 Polynomials = linear regression with X^j basis.
- 2 Step functions = piecewise constant.
- 3 Cubic splines = piecewise cubics joined smoothly.
- 4 Natural splines stabilise the boundaries.
- 5 Smoothing splines: penalised second derivative.
- 6 LOESS: weighted local polynomials.
- 7 GAMs: $\sum_j f_j(X_j)$ — flexible *and* interpretable.

What's next

Chapter 8: *Tree-Based Methods*.

- Decision trees for regression and classification.
- Bagging, random forests, boosting, BART.
- The dominant family of methods for tabular data.


Appendix — what is in here, and why

Nothing in the appendix is needed to follow the chapter: each slide is a formal derivation, a longer worked exercise, or a side topic. Read it when you want the full story, or when a later chapter sends you back here.

Contents of the appendix

- **Wage vs. age: four models side by side** — a second illustration of the cubic-vs-natural contrast already made in the main deck.
- **How the spline basis is built** — one truncated-power term $(x - \xi)_+^3$ drawn on its own, and the constraint counting that turns loose cubic pieces into one smooth curve. The main deck already states the basis and the degrees-of-freedom rule; this is the same argument taken slowly.
- **Regression splines by hand (Extended Exercise 7.1)** — evaluating a spline basis and its fit with pencil and paper. Exercise 7.3 covers the same counting argument in five minutes.


Wage vs. age: four models side by side



Takeaway

All four recover the same rise–plateau–fall shape but differ at the edges: the polynomial dips, the step function is blocky, the natural spline stays smooth and stable.

The building block: one truncated-power term per knot




Reading $(x - \xi)_+^3$

The “+” is the *positive part*: the term is 0 for $x \leq \xi$ and equals $(x - \xi)^3$ afterwards. Its value, slope and curvature are all 0 at ξ , so switching it on adds a bend *only to the right* of the knot without breaking C^2 -smoothness.

Takeaway

Full basis for K knots: $1, x, x^2, x^3, (x - \xi_1)_+^3, \dots, (x - \xi_K)_+^3$ — that is $K + 4$ functions, exactly the $K + 4$ degrees of freedom of a cubic spline.

Adding constraints turns loose pieces into a spline



Takeaway

Top-left: unconstrained pieces jump. Adding continuity of value, then slope, then curvature (bottom-left) removes the kinks — the cubic spline is smoothest.

Extended Exercise 7.1 — Regression splines by hand [Math]

Extended exercise (≈ 15 min) — try it before the solution

Build a cubic regression spline for wage on age with $K = 3$ interior knots $\xi_1 = 25 < \xi_2 = 40 < \xi_3 = 60$.

- 1 Write the truncated-power basis and state the total degrees of freedom.
- 2 The fitted function is

$$\hat{g}(x) = 40 + 0.5x + 0.0004(x - 25)_+^3 - 0.0002(x - 40)_+^3 + 0.0003(x - 60)_+^3.$$

Evaluate $\hat{g}(50)$.

- 3 Show that $\hat{g}, \hat{g}', \hat{g}''$ are continuous at $\xi_2 = 40$ but $\hat{g}'''$ jumps.
- 4 How many degrees of freedom does a *natural* cubic spline with the same three knots use? Explain the difference.

Solution — Extended Exercise 7.1 (1/4)

Solution

(a) **Truncated-power basis.** With $K = 3$ knots the cubic-spline basis is

$$1, x, x^2, x^3, (x - \xi_1)_+^3, (x - \xi_2)_+^3, (x - \xi_3)_+^3,$$

so a fitted spline is $\hat{g}(x) = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3 + \sum_{k=1}^3 \beta_{3+k} (x - \xi_k)_+^3$.

Degrees of freedom. Four base terms plus one per knot: $4 + K = 7$. Equivalently, a free cubic in each of the $K + 1 = 4$ regions has $4 \cdot 4 = 16$ parameters, minus 3 continuity constraints (g, g', g'') at each of the 3 knots: $16 - 9 = 7$.

Why this basis works: each $(x - \xi_k)_+^3$ has value, slope and curvature 0 at its knot, so it bends the fit only to the *right* of ξ_k without breaking smoothness — plain OLS on these 7 columns automatically yields a C^2 piecewise cubic.

Solution — Extended Exercise 7.1 (2/4)

Solution

(b) Evaluate $\hat{g}(50)$. Here the quadratic and cubic *base* coefficients happen to be zero, so only the linear base term and the *active* truncated terms contribute.

The general rule: before any arithmetic, compare x with each knot — a truncated term is switched *on* only where $x > \xi_k$:

Term	$(x - 25)_+^3$	$(x - 40)_+^3$	$(x - 60)_+^3$
At $x = 50$	on ($50 > 25$)	on ($50 > 40$)	off ($50 \leq 60$)

So $\hat{g}(50)$ is a sum of the base terms plus *two* truncated contributions — the third knot has not been reached yet. This bookkeeping step is where the marks are lost; the arithmetic on the next slide is routine.

Solution — Extended Exercise 7.1 (3/4)

Solution

$$(50 - 25)^3 = 25^3 = 15625, \quad (50 - 40)^3 = 10^3 = 1000, \quad (50 - 60)_+^3 = 0.$$

Therefore

$$\begin{aligned} \hat{g}(50) &= 40 + 0.5(50) + 0.0004(15625) - 0.0002(1000) + 0.0003(0) \\ &= 40 + 25 + 6.25 - 0.20 + 0 = 71.05. \end{aligned}$$

Only the basis functions “switched on” at $x = 50$ enter the sum — this locality is the point of the truncated basis.

Common mistake

Evaluating the third term as $(50 - 60)^3 = (-10)^3 = -1000$ instead of 0. The “+” truncation zeroes the term for $x \leq 60$; forgetting it would give the wrong value $71.05 - 0.30 = 70.75$.

Solution — Extended Exercise 7.1 (4/4)

Solution

(c) Continuity at $\xi_2 = 40$. Only the term $t(x) = (x - 40)_+^3$ changes behaviour there. With $t' = 3(x - 40)_+^2$ and $t'' = 6(x - 40)_+$, its one-sided limits at $x = 40$ are

	t	t'	t''
$x \rightarrow 40^-$	0	0	0
$x \rightarrow 40^+$	0	0	0

so $\hat{g}, \hat{g}', \hat{g}''$ are continuous. But $t''' = 6 \cdot 1\{x > 40\}$ jumps from 0 to 6, hence $\hat{g}'''$ is discontinuous — the signature of a cubic spline (continuous curvature, kinked third derivative).

(d) Natural cubic spline. It is forced *linear* beyond the two boundary knots, imposing 2 extra constraints at each end ($g'' = g''' = 0$), i.e. 4 fewer df: $7 - 4 = 3$. A natural spline spends its budget in the interior and is far more stable at the edges.