

Mock Exam 1 — Solutions

Lectures 1–4 (Chapters 1–3)

Prof. Dr. Christoph Weisser

HSBI

Summer Semester 2026

Overview of the mock exam

Problem	Topic	Points
P1	Fundamentals of statistical learning	15
P2	Bias–variance decomposition	15
P3	K -nearest neighbours classification	12
P4	Simple linear regression by hand	24
P5	Multiple regression — interpreting Python output	24
Σ		90

Exam conditions

90 minutes for 90 points (≈ 1 point per minute). Allowed aids: non-programmable calculator and one handwritten A4 sheet (both sides). Justify every answer; round final results to 3 decimals.

Problem 1 — task

Problem 1 (15 P) — Fundamentals of statistical learning

- (a) [4 P] Prediction vs. inference — explain the difference; one example each.
- (b) [4 P] Parametric vs. non-parametric estimation of f ; one advantage of each; which flexibility–interpretability trade-off is behind the choice?
- (c) [3 P] Classify: (i) churn prediction for 800 telecom customers; (ii) which firm characteristics are associated with CEO salary (90 firms); (iii) find groups among 5,000 shopping baskets.
- (d) [4 P] Why is training MSE generally smaller than test MSE? What is overfitting?

Problem 1 — solution (a) and (b)

Solution P1 — (a) and (b)

(a) Prediction: accurate $\hat{Y} = \hat{f}(X)$ for new X ; $\hat{f}$ may be a black box. *Example:* is an incoming e-mail spam?

Inference: understand *how* Y depends on $X_1, \dots, X_p$ (which predictors, direction, size); $\hat{f}$ must be interpretable. *Example:* which ad channels increase sales and by how much?

(b) Parametric: (1) assume a form, e.g. $f(X) = \beta_0 + \sum_j \beta_j X_j$; (2) estimate the few parameters (least squares). *Advantage:* needs less data; simple and interpretable.

Non-parametric: no functional form; $\hat{f}$ follows the data (KNN, splines). *Advantage:* captures many possible shapes of f — no risk of a badly wrong form.

Trade-off: flexibility reduces bias but costs data, risks overfitting and loses interpretability; for inference we usually prefer restrictive models.

Problem 1 — solution (c) and (d)

Solution P1 — (c) and (d)

(c) (i) *Classification*: qualitative response cancel yes/no (prediction).

(ii) *Regression*: quantitative response CEO salary (inference).

(iii) *Unsupervised*: no response variable; goal is clustering.

(d) The method is fitted by minimising *training* error, so $\hat{f}$ adapts to the training observations *including their noise* — training MSE systematically underestimates the error on fresh data. With growing flexibility the training MSE falls monotonically while the test MSE is U-shaped.

Overfitting: the fit follows patterns of the training noise that do not exist in the population; a less flexible model would have *lower test* MSE. Symptom: very small training MSE together with a rising test MSE (widening gap).

Problem 2 — task

Problem 2 (15 P) — Bias–variance decomposition

At a fixed x_0 three methods were analysed; $\sigma^2 = \text{Var}(\varepsilon) = 1.00$:

	A (linear)	B (moderate)	C (very flexible)
$[\text{Bias}]^2$	2.25	0.16	0.01
$\text{Var}(\hat{f}(x_0))$	0.25	1.44	3.60

- (a) **[4 P]** State the decomposition of the expected test MSE and explain the components. (b) **[5 P]** Compute the expected test MSE for A–C; which method to choose? (c) **[3 P]** Smallest expected test MSE any method could achieve here? (d) **[3 P]** Behaviour of bias²/variance/training/test MSE as flexibility grows; U-shape.

Problem 2 — solution (a) and (b)

Solution P2 — (a) and (b)

$$(a) \quad E[(y_0 - \hat{f}(x_0))^2] = \underbrace{\text{Var}(\hat{f}(x_0))}_{\text{sensitivity to training set}} + \underbrace{[\text{Bias}(\hat{f}(x_0))]^2}_{\text{systematic model error}} + \underbrace{\text{Var}(\varepsilon)}_{\text{irreducible noise}}$$

(b) Expected test MSE = Bias² + Var + σ^2 :

	Bias ²	Var	σ^2	Test MSE
A	2.25	0.25	1.00	3.50
B	0.16	1.44	1.00	2.60
C	0.01	3.60	1.00	4.61

Choose **B**: lowest expected test MSE. C pays more in variance (+2.16) than it gains in bias (−0.15).

Problem 2 — solution (c) and (d)

Solution P2 — (c) and (d)

(c) $\sigma^2 = 1.00$. Since $\text{Bias}^2 \geq 0$ and $\text{Var} \geq 0$, $E[(y_0 - \hat{f}(x_0))^2] \geq \text{Var}(\varepsilon) = 1.00$ for every method. Even the true f cannot predict the noise draw ε_0 : the *irreducible error* is the floor.

(d) As flexibility increases:

- Bias^2 *decreases* (typically monotonically); variance *increases*;
- training MSE decreases *monotonically* (the fit hugs the training data);
- test MSE = $\text{Bias}^2 + \text{Var} + \sigma^2$ first falls (bias reduction dominates) then rises (variance dominates) $\Rightarrow$ **U-shape**.

In the table: $3.50 \rightarrow 2.60 \rightarrow 4.61$ for $A \rightarrow B \rightarrow C$ — minimum at intermediate flexibility.

Problem 3 — task

Problem 3 (12 P) — KNN classification

Seven visitors with scores X_1 (pages) and X_2 (minutes) and outcome Y (purchase); new visitor at $x_0 = (4,2)$ — classify with Euclidean distance.

Obs.	1	2	3	4	5	6	7
X_1	4	2	6	4	1	7	5
X_2	3	2	2	5	1	4	6
Y	Yes	No	No	No	Yes	Yes	No

- (a) **[4 P]** All distances; classification with $K = 1$. (b) **[4 P]** Classification with $K = 3$ and $\hat{\text{Pr}}(\text{Yes} \mid x_0)$.
(c) **[2 P]** Effect of K on bias/variance; training error of 1-NN. (d) **[2 P]** Bayes classifier: definition; why unattainable; relation to KNN.

Problem 3 — solution (a) and (b)

Solution P3 — (a) and (b)

(a) $d^2 = (X_1 - 4)^2 + (X_2 - 2)^2$, sorted:

Obs.	1	2	3	4	5	6	7
d^2	1	4	4	9	10	13	17
d	1.000	2.000	2.000	3.000	3.162	3.606	4.123
Y	Yes	No	No	No	Yes	Yes	No

$K = 1$: nearest is Obs. 1 ($d = 1.000$; Yes) $\Rightarrow$ **Yes**.

(b) $K = 3$: neighbours are Obs. 1 (Yes) and Obs. 2 and Obs. 3 (both No). $\widehat{\Pr}(\text{Yes} \mid x_0) = 1/3 \approx 0.333$; $\widehat{\Pr}(\text{No} \mid x_0) = 2/3 \Rightarrow$ majority vote: **No**. Note that $K = 1$ and $K = 3$ disagree — the answer depends on the flexibility.

Problem 3 — solution (c) and (d)

Solution P3 — (c) and (d)

(c) Small K (e.g. 1): highly flexible and jagged boundary $\Rightarrow$ *low bias, high variance*. Large K : strong local averaging, smooth almost linear boundary $\Rightarrow$ *high bias, low variance*. Training error of 1-NN: 0 — each training point is its own nearest neighbour (says nothing about test error).

(d) The *Bayes classifier* assigns x to the class j maximising the *true* conditional probability $\Pr(Y = j \mid X = x)$; it minimises the expected test error (*Bayes error rate* = lower bound for all classifiers). Unattainable in practice: the conditional distribution of Y given X is unknown. KNN *estimates* these probabilities locally from the K nearest neighbours and then classifies to the largest estimate — a data-driven approximation of the Bayes rule.

Problem 4 — task

Problem 4 (24 P) — Simple linear regression by hand

Radio spots x and weekly sales y (in 1,000 EUR) for $n = 5$ weeks; $\sum x_i = 15$; $\sum y_i = 55$; $\sum x_i^2 = 55$; $\sum x_i y_i = 195$.

Week	1	2	3	4	5
x_i	1	2	3	4	5
y_i	6	7	11	13	18

- (a) [6 P] $\hat{\beta}_1$ and $\hat{\beta}_0$; fitted line; interpretation. (b) [4 P] Fitted values and residuals; check $\sum e_i = 0$. (c) [3 P] RSS and RSE; interpretation. (d) [6 P] $SE(\hat{\beta}_1)$; t -test of $H_0: \beta_1 = 0$ ($t_{0.975,3} = 3.182$); 95 % CI. (e) [3 P] TSS and R^2 ; interpretation. (f) [2 P] Prediction at $x = 6$; one caveat.

Problem 4 — solution (a) and (b)

Solution P4 — (a) and (b)

(a) $\bar{x} = 3$; $\bar{y} = 11$; $S_{xy} = 195 - 5 \cdot 3 \cdot 11 = 30$; $S_{xx} = 55 - 5 \cdot 3^2 = 10$

$$\hat{\beta}_1 = \frac{30}{10} = 3.000 \quad \hat{\beta}_0 = 11 - 3 \cdot 3 = 2.000 \quad \hat{y} = 2.000 + 3.000x$$

Slope: +1 spot $\Rightarrow$ expected sales +3,000 EUR per week. *Intercept*: expected sales without spots = 2,000 EUR (mild extrapolation; $x = 0$ not observed).

(b) $\hat{y}_i = 2 + 3x_i$; $e_i = y_i - \hat{y}_i$:

i	1	2	3	4	5
$\hat{y}_i$	5	8	11	14	17
e_i	1	-1	0	-1	1

$\sum e_i = 1 - 1 + 0 - 1 + 1 = 0 \checkmark$ (always true for least squares with intercept).

Problem 4 — solution (c) and (d)

Solution P4 — (c) and (d)

(c) $RSS = 1 + 1 + 0 + 1 + 1 = 4.000$; $RSE = \sqrt{4/(5-2)} = \sqrt{4/3} = 1.155$ (thousand EUR): typical deviation of observed sales from the true line ($\hat{\sigma}$).

$$(d) SE(\hat{\beta}_1) = \frac{RSE}{\sqrt{S_{xx}}} = \frac{1.1547}{\sqrt{10}} = 0.365; \quad t = \frac{3.000}{0.3651} = 8.216 \quad (df = 3).$$

- $|t| = 8.216 > 3.182 = t_{0.975,3} \Rightarrow$ reject $H_0: \beta_1 = 0$ at the 5 % level: significant positive association.
- 95 % CI: $3.000 \pm 3.182 \cdot 0.3651 = 3.000 \pm 1.162 = [1.838; 4.162]$; $0 \notin$ CI — consistent with the test.

Problem 4 — solution (e) and (f)

Solution P4 — (e) and (f)

(e) $TSS = \sum_i (y_i - 11)^2 = 25 + 16 + 0 + 4 + 49 = 94.000$

$$R^2 = 1 - \frac{RSS}{TSS} = 1 - \frac{4}{94} = 0.957$$

95.7 % of the variance of weekly sales is explained by the regression on radio spots (simple regression: $r = \sqrt{R^2} \approx 0.979$).

(f) $\hat{y}(6) = 2 + 3 \cdot 6 = 20.000$ thousand EUR = 20,000 EUR.

Caveat: $x = 6$ lies outside the observed range $[1;5] \Rightarrow$ **extrapolation** — the linear form is unverified there (and $n = 5$ is very small).

Problem 5 — task

Problem 5 (24 P) — Multiple regression output

$n = 60$ employees; salary (1,000 EUR/year); experience (years); female (dummy). OLS fit of salary = $\beta_0 + \beta_1 \text{experience} + \beta_2 \text{female} + \beta_3 (\text{experience} \times \text{female}) + \varepsilon$:

	coef	std err	t	$P > t $
Intercept	46.2000	1.100	42.000	<0.0001
experience	2.3500	0.250	9.400	<0.0001
female	-3.1000	1.550	-2.000	0.0504
experience:female	-0.8500	0.350	-2.429	0.0184

$R^2 = 0.712$; adj. $R^2 = 0.697$; $F = 46.15$ on 3 and 56 df; $\text{Prob}(F) < 0.0001$

(a) [6 P] Interpret all coefficients. (b) [4 P] Significance at 5%; drop female? (c) [4 P] Fitted equation per group. (d) [4 P] Predictions at 10 years; gap. (e) [3 P] R^2 and F -test. (f) [3 P] U-shaped residual plot; VIF = 12.

Problem 5 — solution (a)

Solution P5 — (a) coefficient interpretation

Salary in 1,000 EUR per year; baseline group: men (`female=0`).

- **Intercept 46.20**: expected salary of a *man* with 0 years of experience = 46,200 EUR.
- **experience 2.35**: for *men* each additional year raises expected salary by 2,350 EUR (with an interaction the slope is group-specific).
- **female -3.10**: at 0 years of experience women earn 3,100 EUR *less* on average — the difference of the two intercepts.
- **experience:female -0.85**: the experience slope of women is 0.85 smaller: women gain $2.35 - 0.85 = 1.50$ thousand EUR per year; the expected gender gap *widens* by 850 EUR per year of experience.

Problem 5 — solution (b) and (c)

Solution P5 — (b) and (c)

(b) Intercept and experience: $p < 0.0001 \Rightarrow$ significant. Interaction: consistency $t = -0.85/0.35 = -2.429 \checkmark$; $p = 0.0184 < 0.05 \Rightarrow$ significant. female: $p = 0.0504 > 0.05 \Rightarrow$ *not* significant at exactly 5 % (borderline). Do *not* remove it: **hierarchy principle** — a model containing an interaction should contain the corresponding main effects; and $p = 0.0504$ is weak evidence against an effect, not proof of none.

(c) Group-wise fitted equations:

$$\text{Men: } \widehat{\text{salary}} = 46.20 + 2.35 \text{ experience}$$

$$\text{Women: } \widehat{\text{salary}} = (46.20 - 3.10) + (2.35 - 0.85) \text{ experience} = 43.10 + 1.50 \text{ experience}$$

Different intercepts *and* slopes — exactly what the interaction permits.

Problem 5 — solution (d) and (e)

Solution P5 — (d) and (e)

(d) At 10 years of experience:

$$\text{Man: } 46.20 + 2.35 \cdot 10 = 69.70 \Rightarrow 69,700 \text{ EUR}$$

$$\text{Woman: } 43.10 + 1.50 \cdot 10 = 58.10 \Rightarrow 58,100 \text{ EUR}$$

$$\text{Gap: } 58.10 - 69.70 = -11.60 \text{ thousand EUR} = \hat{\beta}_2 + 10\hat{\beta}_3 = -3.10 - 8.50 \Rightarrow \text{women earn 11,600 EUR less.}$$

(e) $R^2 = 0.712$: the three terms explain 71.2 % of the salary variance. F -test: $H_0: \beta_1 = \beta_2 = \beta_3 = 0$ (no predictor useful) vs. at least one $\beta_j \neq 0$; $F = 46.15$ (3 and 56 df) with $p < 0.0001 \Rightarrow$ reject H_0 : the model is useful overall.

Problem 5 — solution (f) diagnostics

Solution P5 — (f) diagnostics

(i) **U-shaped residual-vs-fitted plot:** systematic curvature the linear model misses $\Rightarrow$ **non-linearity**. Remedies: add a polynomial term (e.g. experience^2) or transform variables (e.g. $\log(\text{salary})$ or $\sqrt{X}$).

(ii) **VIF = 12 for age:** $\text{VIF}_j = 1/(1 - R_{X_j|X_{-j}}^2) = 12 \Rightarrow R_{X_j|X_{-j}}^2 \approx 0.917$: age is almost a linear combination of the other predictors $\Rightarrow$ strong **multicollinearity** (rules of thumb: $\text{VIF} > 5$ or > 10 is problematic).

Consequence: $\text{SE}(\hat{\beta}_{\text{age}})$ inflated by $\sqrt{12} \approx 3.5 \Rightarrow$ unstable coefficient; low power of its t -test; individual effects of collinear predictors cannot be separated — although overall fit and predictions may remain fine.